

```
import pandas as pd
import numpy as np
```

Problem - 1: Perform a classification task with knn from scratch.

1.Load the Dataset: • Read the dataset into a pandas DataFrame

```
df = pd.read_csv('/content/drive/MyDrive/Copy of diabetes_.csv');
df.head(3)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
0	6	148	72	35	0	33.6		0.627	50	1
1	1	85	66	29	0	26.6		0.351	31	0
2	8	183	64	0	0	23.3		0.672	32	1

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
[6] ✓ 0s
```

```
# Finding all datas
print(df.info())

print()## for spaces

# Finding Missing values
print(df.isnull().sum())

print()

# Finding Summary statistics
print(df.describe(include="all"))

print()

# Finding Shape
print(df.shape)

print()

# 6. finding names
print(df.columns)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
... Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768	int64
1	Glucose	768	int64
2	BloodPressure	768	int64
3	SkinThickness	768	int64
4	Insulin	768	int64
5	BMI	768	float64
6	DiabetesPedigreeFunction	768	float64
7	Age	768	int64
8	Outcome	768	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

```
None
```

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction 0
Age              0
Outcome          0
..               ..
```

```

...    Pregnancies      Glucose   BloodPressure  SkinThickness      Insulin \
count  768.000000  768.000000  768.000000  768.000000  768.000000 \
mean   3.845052  120.894531  69.105469  20.536458  79.799479
std    3.369578  31.972618  19.355807  15.952218  115.244002
min    0.000000  0.000000  0.000000  0.000000  0.000000
25%   1.000000  99.000000  62.000000  0.000000  0.000000
50%   3.000000  117.000000  72.000000  23.000000  30.500000
75%   6.000000  140.250000  80.000000  32.000000  127.250000
max   17.000000  199.000000 122.000000  99.000000  846.000000

          BMI  DiabetesPedigreeFunction      Age      Outcome
count  768.000000  768.000000  768.000000  768.000000
mean   31.992578  0.471876  33.240885  0.348958
std    7.884160  0.331329  11.760232  0.476951
min    0.000000  0.078000  21.000000  0.000000
25%   27.300000  0.243750  24.000000  0.000000
50%   32.000000  0.372500  29.000000  0.000000
75%   36.600000  0.626250  41.000000  1.000000
max   67.100000  2.420000  81.000000  1.000000

(768, 9)

Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')

```

2.Handle Missing Data: • Handle any missing values appropriately, either by dropping or imputing them based on the data.

```
[7]  os
  df.isnull().sum()#find missing data
...
...          0
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

```

# Clean column names
df.columns = df.columns.str.strip().str.lower()

# Drop unwanted columns safely
df = df.drop(
    columns=["bp.2s", "bp.2d", "location", "id", "chol", "stab.glu",
             "hdl", "glyhb", "bp.1s", "bp.id", "time.ppn", "age"],
    errors="ignore"
)

# Handle frame column
if "frame" in df.columns:
    df["frame"] = df["frame"].fillna(df["frame"].mode()[0])
    df["frame"] = df["frame"].map({"small": 0, "medium": 1, "large": 2})

# Handle gender column
if "gender" in df.columns:
    df["gender"] = df["gender"].map({"male": 0, "female": 1})

# Fill numeric columns with median
num_cols = ["ratio", "height", "weight", "waist", "hip"]
for col in num_cols:
```

```
▶ # Fill numeric columns with median
num_cols = ["ratio", "height", "weight", "waist", "hip"]
for col in num_cols:
    if col in df.columns:
        df[col] = df[col].fillna(df[col].median())

df.head()
```

... pregnancies glucose bloodpressure skintickness insulin bmi diabetespedigreefunction outcome

0	6	148	72	35	0	33.6	0.627	1
1	1	85	66	29	0	26.6	0.351	0
2	8	183	64	0	0	23.3	0.672	1
3	1	89	66	23	94	28.1	0.167	0
4	0	137	40	35	168	43.1	2.288	1

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
▶ df.isnull().sum()
```

... 0

pregnancies	0
glucose	0
bloodpressure	0
skintickness	0
insulin	0
bmi	0
diabetespedigreefunction	0
outcome	0

dtype: int64

```
▶ import numpy as np
import pandas as pd

# 1 Clean column names (safe habit)
df.columns = (
    df.columns
    .str.strip()
    .str.lower()
    .str.replace(" ", "_")
)

print("Columns in dataset:")
print(df.columns.tolist())

# 2 Separate FEATURES (X) and TARGET (y)
X = df.drop(columns=[ "outcome" ]).values
y = df[ "outcome" ].values

print("x shape:", X.shape)
print("y shape:", y.shape)
```

... Columns in dataset:

```

▶ def train_test_split_scratch(X, y, test_size=0.3, random_seed=42):
    """
    Splits dataset into train and test sets from scratch.
    """

    np.random.seed(random_seed)

    # Generate array of indices
    indices = np.arange(X.shape[0])
    np.random.shuffle(indices)

    # Determine test size
    test_count = int(len(X) * test_size)

    # Split indices
    test_idx = indices[:test_count]
    train_idx = indices[test_count:]

    # Create train and test sets
    X_train = X[train_idx]
    X_test = X[test_idx]
    y_train = y[train_idx]
    y_test = y[test_idx]

    #Printing the results

```

```

#Printing the results
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

return X_train, X_test, y_train, y_test

x_train, x_test, y_train, y_test = train_test_split_scratch(X, y)

...
X_train shape: (538, 0)
X_test shape: (230, 0)
y_train shape: (538,)
y_test shape: (230,)

```

4.Implement KNN:

- Build the KNN algorithm from scratch (no libraries like sickit-learn for KNN).

- Compute distances using Euclidean distance.
- Write functions for:
 - Predicting the class for a single query.
 - Predicting classes for all test samples.
- Evaluate the performance using accuracy.

```

2] 0s
▶ def euclidean_distance(point1, point2):
    """
    Calculate the Euclidean distance between two points in n-dimensional space.

    Arguments:
    point1 : np.ndarray
        The first point as a numpy array.
    point2 : np.ndarray
        The second point as a numpy array.

    Returns:
    float

```

```
    THE EUCLIDEAN DISTANCE BETWEEN THE TWO POINTS.
```

▶ Raises:
ValueError: If the input points do not have the same dimensionality.
"""

```
if point1.shape != point2.shape:
    raise ValueError("Points must have the same dimensions to calculate Euclidean distance.")

distance = np.sqrt(np.sum((point1 - point2) ** 2))
return distance
```

```
euclidean_distance(x[0], x[1])
```

```
... np.float64(0.0)
```

```
▶ try:
```

```
    point1 = np.array([3, 4])
    point2 = np.array([0, 0])

    result = euclidean_distance(point1, point2)

    expected_result = 5.0
    assert np.isclose(result, expected_result), f"Expected {expected_result}, but got {result}"

    print("Test passed successfully!")
except ValueError as ve:
    print(f"ValueError: {ve}")
except AssertionError as ae:
    print(f"AssertionError: {ae}")
except Exception as e:
    print(f"An unexpected error occurred: {e}")

... Test passed successfully!
```

```
▶ x_test_sample = x_test[:5]
y_test_sample = y_test[:5]

predictions = knn_predict(x_test_sample, x_train, y_train, k=3)

print("Predictions:", predictions)
print("Actual labels:", y_test_sample)

assert predictions.shape == y_test_sample.shape, (
    "The shape of predictions does not match the shape of the actual labels."
)

print("Test case passed successfully!")
except AssertionError as ae:
    print(f"AssertionError: {ae}")
except Exception as e:
    print(f"An unexpected error occurred: {e}")

... Predictions: [1 1 1 1 1]
Actual labels: [0 0 0 0 0]
Test case passed successfully!
```

```
def compute_accuracy(y_true, y_pred):
    """ Compute accuracy as a percentage (0 to 100). """
    correct_predictions = np.sum(y_true == y_pred)
    total_predictions = len(y_true)
    accuracy = (correct_predictions / total_predictions) * 100
    return accuracy

try:
    predictions = knn_predict(x_test, x_train, y_train, k=3)

    accuracy = compute_accuracy(y_test, predictions)

    print(f"Accuracy of the KNN model on the test set: {accuracy:.2f}%")
except Exception as e:
    print(f"An unexpected error occurred during prediction or accuracy computation: {e}")

Accuracy of the KNN model on the test set: 34.78%
```

```
for k in [1, 3, 5, 7, 9, 11, 13, 15, 16, 17, 18, 20]:
    preds = knn_predict(x_test, x_train, y_train, k)
    acc = np.mean(preds == y_test)
    print(f"k = {k} -> Accuracy = {acc*100:.2f}%")

...   k = 1 -> Accuracy = 34.78%
      k = 2 -> Accuracy = 34.78%
      k = 3 -> Accuracy = 65.22%
      k = 5 -> Accuracy = 34.78%
      k = 7 -> Accuracy = 34.78%
      k = 9 -> Accuracy = 34.78%
      k = 11 -> Accuracy = 34.78%
      k = 13 -> Accuracy = 34.78%
      k = 15 -> Accuracy = 65.22%
      k = 16 -> Accuracy = 65.22%
      k = 17 -> Accuracy = 65.22%
      k = 18 -> Accuracy = 65.22%
      k = 20 -> Accuracy = 65.22%
```

```
import matplotlib.pyplot as plt
def experiment_knn_k_values(x_train, y_train, x_test, y_test, k_values):
    """
    Run KNN predictions for different values of k and plot the accuracies.

    Arguments:
    x_train : np.ndarray
        The training feature matrix.
    y_train : np.ndarray
        The training labels.
    x_test : np.ndarray
        The test feature matrix.
    y_test : np.ndarray
        The test labels.
    k_values : list of int
        A list of k values to experiment with.

    Returns:
    dict
        A dictionary with k values as keys and their corresponding accuracies as values.
    """
    accuracies = {}

    for k in k_values:
```

```

    predictions = knn_predict(x_test, x_train, y_train, k=k)
    accuracy = compute_accuracy(y_test, predictions)
    accuracies[k] = accuracy

    print(f"Accuracy for k={k}: {accuracy:.2f}%")

plt.figure(figsize=(10, 5))
plt.plot(k_values, list(accuracies.values()), marker='o')
plt.xlabel('k (Number of Neighbors)')
plt.ylabel('Accuracy (%)')
plt.title('Accuracy of KNN with Different Values of k')
plt.grid(True)
plt.show()

return accuracies

k_values = range(1, 21)

try:
    accuracies = experiment_knn_k_values(x_train, y_train, x_test, y_test, k_values)
    print("Experiment completed. Check the plot for the accuracy trend.")

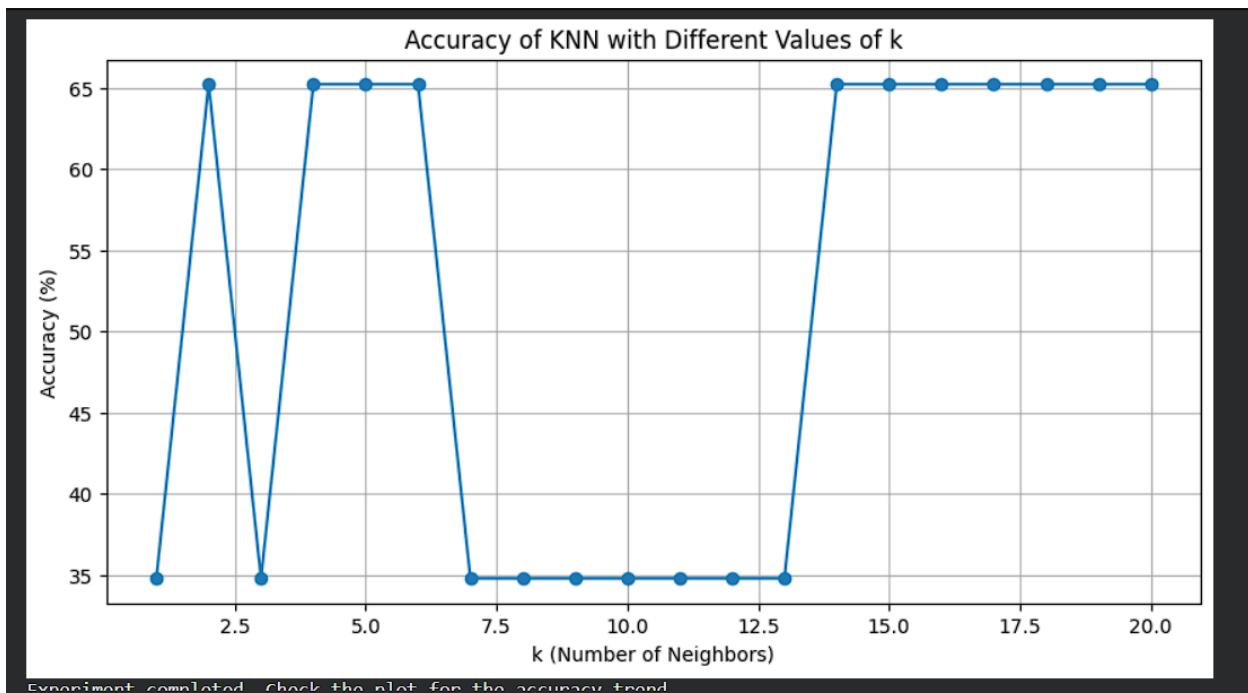
except Exception as e:
    print(f"An unexpected error occurred during the experiment: {e}")

```

```

Accuracy for k=1: 34.78%
Accuracy for k=2: 65.22%
Accuracy for k=3: 34.78%
Accuracy for k=4: 65.22%
Accuracy for k=5: 65.22%
Accuracy for k=6: 65.22%
Accuracy for k=7: 34.78%
Accuracy for k=8: 34.78%
Accuracy for k=9: 34.78%
Accuracy for k=10: 34.78%
Accuracy for k=11: 34.78%
Accuracy for k=12: 34.78%
Accuracy for k=13: 34.78%
Accuracy for k=14: 65.22%
Accuracy for k=15: 65.22%
Accuracy for k=16: 65.22%
Accuracy for k=17: 65.22%
Accuracy for k=18: 65.22%
Accuracy for k=19: 65.22%
Accuracy for k=20: 65.22%

```



```
def min_max_scale(X):
    """
    Manually scale the feature matrix X using min-max scaling.
    Returns the scaled version of X.
    """
    X_min = X.min(axis=0)          # minimum of each column
    X_max = X.max(axis=0)          # maximum of each column

    return (X - X_min) / (X_max - X_min)

x_scaled = min_max_scale(x)
x_scaled

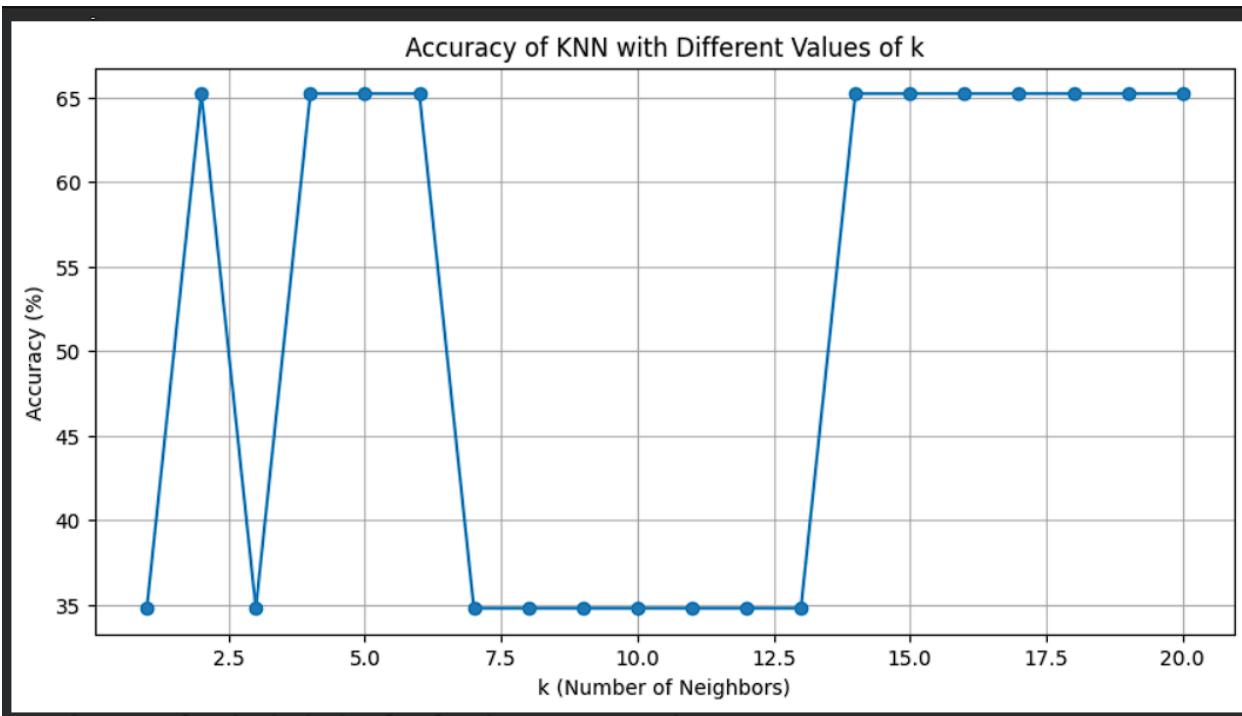
array([1, shape=(768, 0), dtype=float64)
```

```
▶ x_train_s, x_test_s, y_train_s, y_test_s = train_test_split_scratch(x_scaled, y)  
...  X_train shape: (538, 0)  
    X_test shape: (230, 0)  
    y_train shape: (538,)  
    y_test shape: (230,)
```

```
❶ accuracy_scaled = np.mean(y_pred_scaled == y_test_s)
accuracy_scaled
... np.float64(0.34782608695652173)
```

```
❶ k_values = range(1, 21)
try:
    accuracies = experiment_knn_k_values(x_train_s, y_train_s, x_test_s, y_test_s, k_values)
    print("Experiment completed. Check the plot for the accuracy trend.")
except Exception as e:
    print(f"An unexpected error occurred during the experiment: {e}")

... Accuracy for k=1: 34.78%
Accuracy for k=2: 65.22%
Accuracy for k=3: 34.78%
Accuracy for k=4: 65.22%
Accuracy for k=5: 65.22%
Accuracy for k=6: 65.22%
Accuracy for k=7: 34.78%
Accuracy for k=8: 34.78%
Accuracy for k=9: 34.78%
Accuracy for k=10: 34.78%
Accuracy for k=11: 34.78%
Accuracy for k=12: 34.78%
Accuracy for k=13: 34.78%
Accuracy for k=14: 65.22%
Accuracy for k=15: 65.22%
Accuracy for k=16: 65.22%
Accuracy for k=17: 65.22%
Accuracy for k=18: 65.22%
Accuracy for k=19: 65.22%
```



```

▶ import time
k_values = range(1, 16)

scaled_accuracy = []
unscaled_accuracy = []
time_unscaled = []
time_scaled = []

for kk in k_values:
    # Unscaled
    start = time.time()
    unscaled_pred = knn_predict(x_test, x_train, y_train, kk)
    end = time.time()

    unscaled_accuracy.append(compute_accuracy(y_test, unscaled_pred))
    time_unscaled.append(end - start)

    # Scaled
    start = time.time()
    scaled_pred = knn_predict(x_test_s, x_train_s, y_train, kk)
    end = time.time()

    scaled_accuracy.append(compute_accuracy(y_test, scaled_pred))
    time_scaled.append(end - start)

plt.plot(k_values, unscaled_accuracy, marker='o', label='Unscaled', color='red')
plt.plot(k_values, scaled_accuracy, marker='o', label='Scaled')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('k vs Accuracy')
plt.legend()
plt.grid(True)
plt.show()

```

k	Unscaled Accuracy	Scaled Accuracy
1	~52%	~52%
2	~55%	~55%
3	~58%	~58%
4	~60%	~60%
5	~62%	~62%
6	~64%	~64%
7	~65%	~65%
8	~65%	~65%
9	~65%	~65%
10	~65%	~65%
11	~65%	~65%
12	~65%	~65%
13	~65%	~65%
14	~65%	~65%
15	~65%	~65%
16	~65%	~65%

The accuracy of the KNN model is low and unstable for small values of k because the model is sensitive to noise. As k increases, accuracy improves since predictions are based on more neighbors, making them more stable. The scaled dataset consistently performs better than the unscaled dataset because scaling ensures fair distance calculations. The prediction time remains almost constant for all values of k, so increasing k does not significantly increase computational cost. Based on the plots, the optimal value of k lies around 13–15, where the scaled data achieves the highest accuracy with reasonable computation time.

Problem - 4 - Additional Questions {Optional - But Highly Recommended}: • Discuss the challenges of using KNN for large datasets and high-dimensional data. • Suggest strategies to improve the efficiency of KNN (e.g., approximate nearest neighbors, dimensionality reduction).

Challenges of using KNN for large datasets and high-dimensional data: KNN becomes computationally expensive for large datasets because it calculates distances between every test

sample and all training samples. It also consumes a lot of memory as it stores all training data. For high-dimensional data, the “curse of dimensionality” makes distances less meaningful, causing predictions to become less accurate. Additionally, KNN is sensitive to irrelevant features and feature scales, which can further reduce its performance.

Strategies to improve KNN efficiency: Efficiency can be improved by using approximate nearest neighbor methods like KD-Trees or Ball Trees, which reduce the number of distance calculations. Dimensionality reduction techniques such as PCA can help simplify the feature space, making distances more meaningful and faster to compute. Feature scaling ensures no single feature dominates the distance metric, and selecting the optimal K through cross-validation improves accuracy. Clustering can also be used to limit search space for faster predictions.